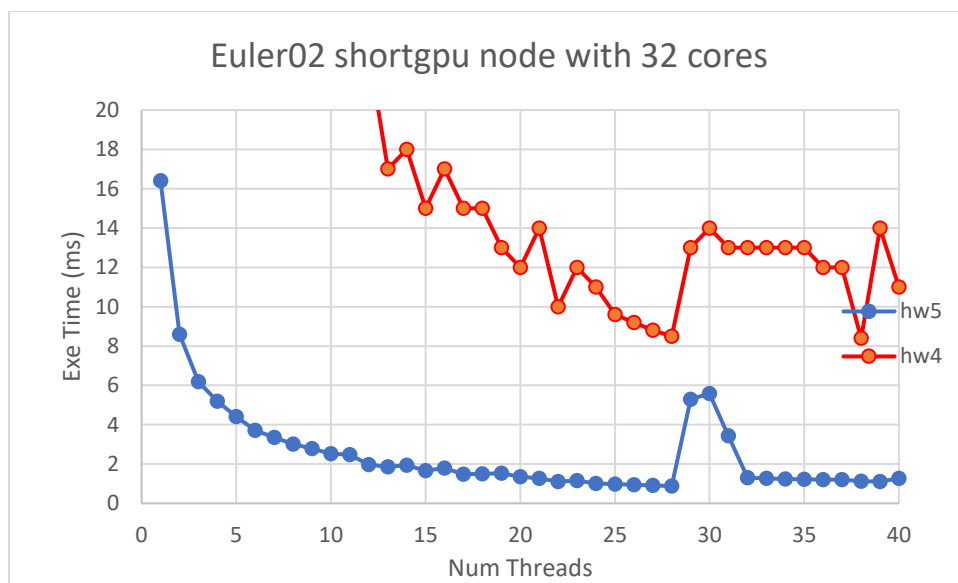


Zoomed in plot:



I used the shortgpu node with $N=32$. There was significant improvement when using vectorization over my code. Using 1 thread the vectorized program was able to compute the result almost as fast as the non-vectorized code from homework 4 using 20 threads.